

DBMS ANAND BIHARI : MOD1 PDF2

Page 1 — “Database Management System” (Title)

What the page is doing: It sets the scope: you’re entering DBMS—software that lets you **define, create, maintain, and control** databases.

Deep explanation: In exam answers, a tight definition is:

A DBMS is software that provides systematic ways to store data, run queries, ensure correctness (constraints/transactions), protect access (authorization), and keep data safe (backup/recovery). Connect this to why we don’t just use files: we need **concurrency control, data integrity, abstraction, security, and performance**.

Mini example: Consider a college portal. A DBMS avoids two admins overwriting each other’s updates (transactions + locking) and ensures rollbacks if a fee-payment step fails.

Exam tip: Start long answers with definition → problems with file systems → benefits of DBMS (integrity, independence, concurrency, recovery, security).

Page 2 — “Data Models”

Slide idea: What a **data model** is and why we need it. The slide says it’s a collection of concepts to describe database structure; often shown pictorially; shows relationships; and mentions **conceptual, logical, physical** levels; also PK/FK.

Deep explanation:

A **data model** is a *vocabulary plus rules* for describing data. It decides the building blocks (tables/entities, attributes, relationships) and how integrity is expressed (keys/constraints). The three classic **levels** separate *what the data means* from *how it’s stored*:

- **Conceptual:** business view (entities like Student, Course, and their relationships).
- **Logical:** DBMS-agnostic “table/columns/keys” detail suitable for implementation on any RDBMS.
- **Physical:** file organizations, indexes, partitions—how bytes are laid out.

Why the split matters: It enables **data independence**—you can change storage decisions (e.g., add an index) without changing application code that uses the logical view.

Common pitfall: Students mix up “conceptual vs logical.” **Heuristic:**

Rectangles/diamonds/ovals → conceptual ER; *tables/PK/FK* → logical schema.

Page 3 — “Type of Data Model”

Slide bullets: Conceptual (what the system contains), Logical (how to implement), Physical (how on DBMS), and **Representational/Implementation** (e.g., relational).

Deep explanation (+ examples):

- **Conceptual (ER/Ontology level):** “Student *enrolls in* Course.” Attributes like RollNo, Name.
- **Logical (Relational):** STUDENT(roll_no PK, name, ...), COURSE(course_id PK, ...), ENROLL(roll_no FK→STUDENT, course_id FK→COURSE, PK(roll_no, course_id)).
- **Physical:** ENROLL clustered by course_id; B-tree index on roll_no.
- **Representational:** In practice, we use **relational model** (tables/relations, tuples, attributes) in most commercial systems.

Exam tip: MCQs love asking which level handles **indexes** (physical) vs **PK/FK** (logical) vs **entities/relationships** (conceptual).

Page 4 — “Conceptual Model”

Slide bullets: Conceptual model captures highest-level relationships; 3 elements: **Entity**, **Attribute**, **Relationship**.

Deep explanation:

- **Entity:** a distinguishable object set (STUDENT, COURSE).
- **Attribute:** properties (Student.name, Course.credits). Identify **key attributes** early (e.g., roll_no, course_id).
- **Relationship:** semantic association (ENROLLS_IN). Cardinalities (1–1, 1–N, M–N) and participation (total/partial) encode business rules before you even think tables.

Worked micro-example (ER thinking):

“Each course is offered by exactly one department; each department offers many courses.” → Relationship OFFERED_BY with (COURSE N:1 DEPARTMENT).

Pitfall: Jumping to tables too soon. First, ensure conceptual rules capture reality.

Page 5 — (ER diagram visuals likely)

Slide likely shows ER figures (entities/relationships). The gist is *notation literacy*.

Deep explanation of shapes & use:

- **Rectangle** = entity; **double rectangle** = weak entity.
- **Oval** = attribute; **double oval** = multivalued attribute; **dashed oval** = derived.
- **Diamond** = relationship (labeled with verb); edge multiplicity shows cardinality.

Exam move: If you get a word problem, jot **entities** → **attributes** → **relationships** → **cardinalities** in that order, then draw. (Even if the exam is theory, this structure clarifies your written explanation.)

Page 6 — “Logical Data Model”

Slide bullets: Defines data structure and relationships in detail, independent of physical storage. Features: includes all entities/relationships; all attributes; **PK specified; FK specified; normalization occurs here.**

Deep explanation:

Moving from ER to relational:

- Each **entity** → a **table** with a **primary key**.
- Each **1–N relationship** → put the **FK** on the N-side table.
- Each **M–N relationship** → create a **bridge (associative) table** with composite PK of the two FKs.
- Ensure **normalization** (1NF: atomic columns; 2NF: no partial dependency on a composite key; 3NF: no transitive dependency on non-keys).

Example: ENROLL(roll_no, course_id, semester, grade) with PK(roll_no, course_id), FKs to STUDENT and COURSE.

Pitfall: Keeping multivalued attributes as repeated columns. Correct approach: **separate table**.

Page 7 — “Steps for Designing the Logical Data Model”

Slide bullets: Specify PKs; find relationships; list all attributes; resolve **M–M**; normalize.

Deep explanation with micro-walkthrough:

1. **PKs:** Choose minimal, stable identifiers (avoid name/phone).
2. **Relationships:** Tag each with cardinality; choose FK placement (1–N) or a new table (M–N).

3. **Attributes:** Include domains (e.g., credits INT CHECK (credits BETWEEN 1 AND 5)).
4. **Resolve M–N:** ENROLL as above.
5. **Normalize:** Move derived attributes out (e.g., totalCredits is computed), remove duplicates, enforce dependencies.

Exam tip: If asked “how to convert ER→Relational,” list exactly these steps in order, with one short example.

Page 8 — (Worked example continuation)

What to do here: If the slide shows a conversion example, your answer should **reconstruct the reasoning**:

- Identify each entity and its PK decision (e.g., COURSE uses course_id).
- Show how a relationship became an FK (1–N) or a new table (M–N).
- State **why** (cardinality + dependency).
- Show any **normalization** fix you applied (e.g., moved instructor_phone out of COURSE to an INSTRUCTOR table).

Pro-tip: Always annotate **constraints**: UNIQUE, NOT NULL, CHECK, and **referential actions** (ON DELETE CASCADE/SET NULL)—these are logical-level integrity rules that earn marks.

Page 9 — “Physical Data Model”

Slide bullets: How the model is built in DB; shows **table structures, column names, data types, constraints, PK, FK, relationships**.

Deep explanation:

This is **implementation detail** in a target DBMS (say, PostgreSQL). Decisions include:

- **Data types** (VARCHAR vs TEXT; NUMERIC vs INT).
- **Indexes** (B-tree for equality/range; maybe composite indexes).
- **Partitioning** (by semester or department).
- **Storage** (heap vs clustered index; fillfactor).
- **Physical constraints** and **referential actions**.

Example DDL fragment (illustrative):

```
CREATE TABLE ENROLL(roll_no INT REFERENCES STUDENT(roll_no) ON DELETE CASCADE,
course_id TEXT REFERENCES COURSE(course_id), semester CHAR(5),
grade CHAR(2), PRIMARY KEY(roll_no, course_id));
```

Pitfall: Creating indexes blindly. Tie index choice to **query workload** (e.g., lookups by course_id + semester → composite index (course_id, semester)).

Page 10 — (Physical-level illustration)

What to explain: Likely shows a sample schema diagram or DDL excerpt. Translate diagrams into **concrete choices**:

- **Naming conventions** (snake_case vs camelCase) → consistency improves maintenance.
- **NULL policy** (e.g., grade might be NULL until finalized).
- **Check constraints** (credits in 1..5; semester like 2025s).
- **Index strategy** (FK columns usually indexed to speed joins).
- **Storage optimization** (e.g., BOOLEAN flags vs CHAR(1); avoid oversized types).

Extra exam nugget: Distinguish **logical** (ALTER TABLE ADD COLUMN) vs **physical** tuning (CREATE INDEX, VACUUM/ANALYZE, partitioning).

Page 11 — “Data Model” (Overview List)

Slide bullets: Hierarchical, Network, ER Model, Relational Model, Object-based (OO & Object-relational), Semi-structured (XML).

Deep explanation:

These are **different families of data models** based on structure and rules:

1. **Hierarchical:**
 - Tree-like (root → children).
 - **1:N** relationships only (each child has exactly one parent).
 - Data is accessed by navigating from the root down.
 - Example: IBM IMS DB — “Company” has “Departments” which have “Employees.”
2. **Network:**
 - Graph structure (record types + set types).
 - Supports **M:N** directly.
 - Example: CODASYL DBTG model.
3. **ER Model:**
 - Conceptual-level representation for designing before implementation.
4. **Relational Model:**
 - Tables (relations) with rows (tuples) and columns (attributes).
 - Most popular in commercial DBMS.
5. **Object-based Models:**
 - Combine database + object-oriented programming concepts (classes, inheritance, methods).
 - Object-relational: extends relational with object types.
6. **Semi-structured:**
 - Data doesn’t have a fixed schema; tags or markers indicate structure.
 - Example: XML, JSON databases.

Exam tip: They often ask: *Which model supports M:N directly without extra tables?* → **Network Model.**

Page 12 — “Hierarchical Model”

Slide bullets: Tree-like, child has only one parent, 1:N relationship.

Deep explanation:

- **Structure:** Think **org chart**. Each parent can have many children, but each child belongs to exactly one parent.
- Stored physically as **linked records**: parent node pointers lead to child nodes.
- Navigation: start at root and traverse. No arbitrary queries without navigation.

Example:

```
Company
├── Department (HR, IT)
│   └── Employee
```

Advantages:

- Fast access if you always know the path.
- Good for fixed, stable relationships.

Limitations:

- Poor at representing many-to-many.
 - Changes in structure require rewriting access code.
-

Page 13 — “Network Model”

Slide bullets: Graph-like; child can have multiple parents; supports many-to-many; uses “records” and “sets.”

Deep explanation:

- Records are connected via **pointers** in “sets” (owner → member).
- You can navigate via multiple paths — flexible but more complex.

Example:

A project can have multiple employees; an employee can work on multiple projects — in network DBMS, these are linked without a separate associative table.

Pros:

- Efficient for complex M:N relations.

Cons:

- Application code must manage navigation; not as declarative as SQL.
-

Page 14 — “Entities Relationship Model”

Explanation:

This is the conceptual design standard. Already covered in **Page 4**, but here the emphasis is:

- ER diagrams = blueprint before converting to logical model.
- Includes entity types, attributes, relationships, and constraints.
- Can be extended to **Enhanced ER (EER)** for specialization/generalization.

Extra exam point: Be ready to draw ER diagrams with correct notation for derived/multivalued attributes.

Page 15 & 16 — “Relational Model”

Slide points:

- Data is in **relations** (tables).
- Each row = tuple, each column = attribute.
- Keys ensure uniqueness; FKs ensure referential integrity.

Deep explanation:

- **Domain:** allowed values for an attribute.
- **Relation schema:** name + attributes + domains.
- **Relation instance:** actual rows at a point in time.
- Operations: SELECT, PROJECT, JOIN, etc.

Example schema:

```
Student(roll_no INT PRIMARY KEY, name VARCHAR, dept_code CHAR(4) REFERENCES
Department(dept_code))
```

Exam trap: A relation \neq table in a loose sense — in theory, order of tuples and attributes doesn't matter; duplicates aren't allowed.

Page 17 — “Object-Based Data Model”

Slide points: OO + DB concepts.

Explanation:

- Stores complex data as **objects** with attributes + methods.
 - Supports **encapsulation**, **inheritance**, **polymorphism** inside DB.
 - **Object-Relational** DBMS: relational + user-defined types, table inheritance.
 - Example: PostgreSQL supports custom types and table inheritance.
-

Page 18 — “Semi-Structured Data Model”

Explanation:

- Schema not strictly fixed; data can vary in structure.
- XML and JSON store both data and metadata in the same file.
- Useful for web data, sensor data, or any scenario where schema evolves.

Example:

```
<Student>
  <Name>John</Name>
  <Email>john@example.com</Email>
  <Hobbies>
    <Hobby>Chess</Hobby>
    <Hobby>Reading</Hobby>
  </Hobbies>
</Student>
```

Page 19 — “Schema”

Slide points: Logical structure of DB; analogous to type info; physical schema (storage); logical schema (logical design). Instances = actual data at a time.

Deep explanation:

- **Schema:** *meta-definition* — structure and constraints of data.
 - **Physical schema:** describes files, indexes.
 - **Logical schema:** describes tables, keys, constraints.
 - **External schema:** views for specific user groups.
 - **Instance:** the “snapshot” of the data right now.
 - **Database state:** empty, initial, current.
-

Page 20 — “Instances & States”

Deep explanation:

- **Empty state:** schema created but no data yet.
- **Initial state:** loaded with initial data.
- **Current state:** as of now, possibly changed due to transactions.

Exam tip: They sometimes ask: *Can schema change often?* → Rarely. Instances change frequently.

Page 21 — Schema Recap / Transition to Architecture

This slide likely serves as a bridge from “schema” to “three-schema architecture.”

Key point to reinforce:

- Schema levels are **layers of abstraction**.
- Need for **abstraction**:
 - Applications should not break if the storage method changes.
 - Different users should see customized views.
 - We need an organized way to separate *physical concerns* from *logical concerns*.

Exam link: This is the foundation for the *Three-Schema Architecture* questions.

Page 22 — Three-Schema Architecture: Intro

Slide bullets:

- Three characteristics of DBMS: self-describing, program–data insulation, support multiple views.
- Objective: separate user apps from the physical database.
- Three levels = three schemas = three types of abstraction.

Deep explanation:

Why we need it:

- Without this, changing the storage format (e.g., CSV → binary file) would require rewriting every app that queries the DB.
- Users in different departments need different slices/views of the same DB.

The 3 levels:

1. **Internal Level:** Physical storage details (how data is stored).
2. **Conceptual Level:** Entire logical structure for the community of users.

3. **External Level:** Individual user views (subsets or formatted versions of conceptual data).

Analogy:

Think of a library:

- **Internal:** How books are arranged on shelves.
 - **Conceptual:** Catalog describing all books.
 - **External:** Your personal reading list from the catalog.
-

Page 23 — Architecture Diagram

Likely content: A diagram showing the three layers, with arrows for mappings.

Explanation of the diagram:

- Top: External schemas (many possible).
 - Middle: One conceptual schema.
 - Bottom: One internal schema.
 - **Mappings:**
 - External ↔ Conceptual mapping: transforms user requests into conceptual terms.
 - Conceptual ↔ Internal mapping: transforms logical requests into storage-level actions.
-

Page 24 — Internal Schema

Slide bullets: Describes physical storage structures & access paths (indexes, files). Typically uses a physical data model.

Deep explanation:

- Examples: heap files, sorted files, hashed files.
- Index examples: B+ trees for range search; hash indexes for equality search.
- Access paths: ways to retrieve data efficiently without scanning the whole file.

Exam example:

Q: Which schema level specifies whether an index is B+ tree or hash? → **Internal schema.**

Page 25 — Conceptual & External Schemas

Conceptual schema:

- Describes the *entire* database logically — tables, attributes, constraints, relationships.
- Independent of physical details.
- Example: ER diagram converted to relational schema.

External schemas:

- Define what a *particular* user/application sees.
- Example: “Accounts Department View” with Salary and Payroll tables, hiding academic details.

Exam point: External schemas often correspond to **SQL Views**.

Page 26 — Data Independence: Definition

Slide bullets:

- Logical data independence: change conceptual schema without changing external schemas.
- Physical data independence: change internal schema without changing conceptual schema.

Deep explanation:

Logical DI:

- You add a new table or column at the conceptual level; old user views keep working.

Physical DI:

- You change from heap to indexed storage; logical schema stays the same.

Analogy:

Logical DI is like rearranging chapters in a book without changing the table of contents for readers.

Physical DI is like printing the same book on different quality paper without changing the contents.

Page 27 — Data Independence: Effects

Slide bullets:

- If schema at lower level changes, only mappings need to change.
- Higher-level schemas remain unchanged.
- Apps refer to external schemas, so no need to modify code if DI is preserved.

Example:

If a table is split into two internally for optimization, as long as the conceptual schema still shows it as one table, applications are unaffected.

Page 28 — (Likely) Summary of DI & Architecture

This page probably summarizes:

- **Benefit:** Isolation between layers → easier maintenance, less code breakage.
- **Reality:** Full logical DI is hard; physical DI is easier to achieve in modern DBMS.

Exam tip: MCQs often ask “Which is harder to achieve: logical or physical DI?” → Logical.

Page 29 — Database System Environment

Slide bullets: DBMS software + OS + compiler of programming language = Database system environment.

Deep explanation:

Think of the “ecosystem” needed for DBMS to work:

- **Hardware:** Servers, storage.
- **OS:** Manages resources for DBMS processes.
- **DBMS software:** Query processor, storage manager, transaction manager, etc.
- **Language compiler:** If your application embeds SQL (e.g., C with embedded SQL), compiler + pre-compiler process it.

Exam point: This is the “stack” — changes in OS or hardware can affect DBMS performance.

Page 30 — (Likely) Component Overview Transition

This page likely transitions to the computational model or DBMS components.

Exam value: Be prepared to describe **how** DBMS fits into the broader computing environment — OS schedules DB processes, DBMS manages data logic, apps send SQL requests through APIs.

Page 31 — Computational Model of DBMS

What it likely shows: A block diagram of DBMS components interacting with applications, the OS, and storage.

Deep explanation:

- **Application Program** → sends SQL or embedded queries to the DBMS.

- **DBMS** has multiple modules: parser, optimizer, execution engine, transaction manager, storage manager.
- **OS** provides memory management, process scheduling, and file system access.
- **Storage** holds database files and indexes.

Exam example:

Q: In the computational model, where is query optimization done? → Inside the **DBMS query processor** before execution.

Page 32 — Components of DBMS Software (Part 1)

Slide bullets:

- **DDL Compiler:** Processes schema definitions; stores metadata in the catalog.
- **Query Compiler:** Parses, checks syntax & names, converts query to internal form.
- **Query Optimizer:** Removes redundancy, rearranges operations, generates best execution plan.

Deep explanation:

1. **DDL Compiler**
 - Example: `CREATE TABLE Student(...)` → stored in catalog (data dictionary).
 - Catalog contains schema metadata (table names, attributes, types, constraints).
2. **Query Compiler**
 - Breaks down SQL into tokens, checks if tables/columns exist, types match.
 - Converts to internal representation (query tree).
3. **Query Optimizer**
 - Chooses the cheapest execution plan (cost-based optimization).
 - Example: Decides whether to use an index or full table scan.

Exam tip: If they ask, “Where is the decision to use a join method made?” → **Query optimizer**.

Page 33 — Components of DBMS Software (Part 2)

Slide bullets:

- **Pre-compiler:** Extracts embedded SQL (DML) from host language code.
- **DML Compiler:** Compiles DML into low-level code.
- **Runtime Database Processor:** Executes commands & queries; handles transactions.
- **Stored Data Manager:** Controls access to data and metadata.

Deep explanation:

- **Pre-compiler:** If you write C code with SQL inside (`EXEC SQL ...`), the pre-compiler pulls out SQL for the DBMS compiler to handle.
 - **DML Compiler:** Converts data manipulation commands (`INSERT`, `UPDATE`, `DELETE`) into execution steps.
 - **Runtime DB Processor:** Executes the execution plan, ensures transaction rules (ACID) are met, interacts with concurrency and recovery managers.
 - **Stored Data Manager:** The lowest-level module in DBMS — requests pages from storage, manages buffer pool, enforces security.
-

Page 34 — Database System Utilities (Part 1)

Slide bullets:

- **Loading utility:** Loads external data into the DB.
- **Conversion tool:** Transfers data from one DBMS to another.
- **Backup utility:** Creates full or incremental backups.

Deep explanation:

- **Loading:** Reads raw data (CSV, text, binary) and inserts it into tables.
- **Conversion:** Handles format changes when migrating (e.g., MySQL → PostgreSQL).
- **Backup:**
 - *Full:* copy all data.
 - *Incremental:* only changes since last backup.

Exam tip: If they ask about restoring after failure, mention both **backup utility** and **recovery manager**.

Page 35 — Database System Utilities (Part 2)

Slide bullets:

- **Storage reorganization:** Rearranges how data is stored for better performance.
- **Performance monitoring tool:** Tracks DB usage, gives stats for tuning.

Deep explanation:

- **Reorganization:** E.g., rebuilding indexes, clustering tables, defragmenting files.
 - **Monitoring:** Tracks slow queries, deadlocks, disk I/O stats. Helps DBAs decide if indexes or partitioning are needed.
-

Page 36 — Database System Utilities (Part 3)

Slide bullets:

- **Sorting files, data compression, monitoring access by users.**

Deep explanation:

- Sorting is used in query execution (ORDER BY, GROUP BY).
- Compression reduces storage use, can speed up reads (less I/O).
- User access monitoring helps detect misuse or security breaches.

Exam example:

Q: Which DBMS utility would help in detecting suspicious query patterns? → **User access monitoring tool.**

Page 37 — DBMS Architecture

Slide bullets:

- **1-tier:** Direct DB access.
- **2-tier:** Client ↔ Server (via JDBC/ODBC).
- **3-tier:** Middle layer handles business logic.
- **N-tier:** More layers (presentation, business, data).

Deep explanation:

- **1-tier:** Developer works directly on DBMS — good for testing, not for production.
 - **2-tier:** Fat client connects directly to DB server.
 - **3-tier:** Client talks to application server, which talks to DB server — improves security & scalability.
 - **N-tier:** Used in large distributed systems.
-

Page 38 — 2-tier and 3-tier Architecture Diagram

Explanation:

- 2-tier: Client → DB Server directly.
- 3-tier: Client → App Server → DB Server.
- Middle tier in 3-tier handles:
 - Transaction management
 - Business rules
 - Security enforcement

Exam trap: If they ask “Which architecture is most secure?” → 3-tier (middle tier shields DB from direct access).

Page 39 — Centralized & Client-Server DBMS

Slide bullets:

- **Centralized:** All processing done at one site; users connect via terminals.
- **Client-Server:** Specialized servers for DB, web, email, files, etc.

Deep explanation:

- Centralized systems are simpler but can be a bottleneck.
 - Client-server splits processing between clients and dedicated servers, allowing better performance and modularity.
-

Page 40 — Physical Centralized Architecture

Explanation:

- One physical machine runs DBMS, OS, application logic, and handles all storage.
- Clients are “dumb terminals” — no local processing.

Exam tip: They might compare **centralized** vs **distributed** — remember, centralized = one machine; distributed = multiple connected sites.

Page 41 — Basic 2-Tier Client-Server Architectures

Slide bullets:

- Specialized servers: print server, file server, DBMS server, web server, email server.
- Clients can access specialized servers as needed.

Deep explanation:

- In **basic 2-tier**, the client connects directly to a specific server type for a specific function.
- For DBMS: client app uses **ODBC/JDBC** to connect to a **DB server** that executes queries and returns results.
- Separation improves efficiency — the DB server focuses only on database tasks.
- Example: A desktop accounting app (client) talking directly to a MySQL DB server.

Exam trap: They might ask: “In 2-tier, where does SQL execution happen?” → On the **server**, not the client.

Page 42 — Logical Two-Tier Client-Server Architecture

Deep explanation:

- **Client side:** Presentation logic (UI), some application logic, and API calls (ODBC/JDBC).
- **Server side:** Database logic, query processing, transaction handling.
- Communication over a network protocol (e.g., TCP/IP).

Advantage: Less server load if client does some processing.

Disadvantage: Harder to maintain — updates must be deployed to each client.

Page 43 — Client

Slide bullets:

- Provides interface to access & use server resources.
- Can be diskless or have local storage.
- Connected via LAN, wireless, etc.

Deep explanation:

- **Thin client:** Minimal local processing, mainly for presentation.
- **Thick (fat) client:** Significant processing power, can run business logic locally.
- Example: Browser (thin client) vs installed enterprise app (thick client).

Exam example:

Q: Which client type is easier to maintain in large organizations? → Thin client.

Page 44 — DBMS Server

Slide bullets:

- Provides query & transaction services.
- Often called SQL server, query server, or transaction server.
- Accessed via standard APIs: ODBC, JDBC.

Deep explanation:

- **SQL Server:** Executes SQL queries, returns results.
- **Transaction server:** Ensures ACID compliance across multiple operations.
- Needs **client driver** for API used (e.g., MySQL ODBC driver).
- The server may also manage stored procedures, triggers, and indexing.

Page 45 — Two-Tier Client-Server Architecture

Deep explanation:

- Clients can connect to **multiple data sources** (not just one DB).
- Data source can be:
 - RDBMS
 - File-based DB
 - Another data service (e.g., XML API)
- In some designs, client also handles optimization and recovery (seen in object DBMS).

Exam tip: They might ask, “In 2-tier, can one client connect to multiple DBs?” → Yes, via multiple connections/drivers.

Page 46 — Three-Tier Client-Server Architecture

Slide bullets:

- Common for web applications.
- Intermediate layer (application server/web server) between client & DB server.
- Can enhance security: DB server only accessible via middle tier.

Deep explanation:

- **Client:** UI and presentation logic (browser, mobile app).
- **Application server:** Processes business logic, validates requests, formats data.
- **Database server:** Executes SQL, manages transactions.
- Security benefit: Clients cannot directly run arbitrary SQL on DB.

Example: Online banking — browser (client) → bank’s app server → DB server.

Page 47 — Three-Tier Client-Server Architecture Diagram

Deep explanation:

- Shows physical separation between **presentation**, **business logic**, and **data** layers.
- Often implemented as:
 - **Presentation:** HTML/CSS/JS in browser.
 - **Application logic:** Java/PHP/.NET server.
 - **Data:** MySQL/Oracle/PostgreSQL DB.

Exam trap: If asked “Which tier handles input validation?” — usually the **middle tier**.

Page 48 — Classification of DBMS: By Data Model & Users

Slide bullets:

1. By data model: RDBMS, OODBMS, hierarchical, network.
2. By users: Single-user vs multi-user.

Deep explanation:

- **Single-user:** One person at a time (e.g., MS Access).
- **Multi-user:** Supports concurrent access with concurrency control (e.g., MySQL, Oracle).

Exam tip: In multi-user systems, be ready to explain **locks** and **isolation levels**.

Page 49 — Classification of DBMS: By Distribution

Slide bullets:

- **Centralized:** All in one site.
- **Distributed:** Data & DBMS across multiple sites.
 - Homogeneous: same DBMS everywhere.
 - Heterogeneous: different DBMS at different sites.
- **Federated:** Distributed + heterogeneity + local autonomy.

Deep explanation:

- Homogeneous DDBMS is easier to manage — same query language and DBMS version everywhere.
 - Federated DBMS allows each site to keep its own DBMS and control over local data.
-

Page 50 — Classification of DBMS: By Purpose & Cost

Slide bullets:

- **Special-purpose:** Built for one function (airline reservations).
- **General-purpose:** MySQL, Oracle, PostgreSQL.
- **Cost:** From free open-source to multi-million enterprise systems.

Deep explanation:

- Licensing models:
 - Site license (unlimited users at one site).
 - Seat license (limit on concurrent users).
 - Single-user license.
- Enterprise DBMS often modular: replication, partitioning, analytics.

Exam point: If given examples, be able to classify DBMS into these categories.

Page 51 — Classification of DBMS Based on Cost (Expanded Details)

Slide bullets:

- Giant systems sold in modular form: distribution, replication, parallel processing, mobile support.
- Many configurable parameters.
- Licensing models:
 - **Site license** — unlimited copies at one location.
 - **Concurrent user license** — limits simultaneous users.
 - **Standalone single-user** — e.g., MS Access.

Deep explanation:

- **Modular design** allows organizations to buy only what they need (e.g., pay extra for analytics or spatial data).
- **Enterprise DBMS** require careful tuning of parameters: memory allocation, cache size, buffer pool configuration.
- Licensing has major cost implications — for big companies, per-CPU or per-core licensing can be more expensive than per-user.

Exam tip: Be prepared for MCQs that ask: *Which license type allows unlimited concurrent users at a site?* → **Site license**.

Page 52 — Cost Considerations for DBMSs

Slide bullets:

- Cost range: free open-source → millions for enterprise setups.
- Examples of free: MySQL, PostgreSQL.
- Commercial DBMS offer extra modules (time-series, spatial, XML).
- Extra modules sometimes called **cartridges** (Oracle) or **blades**.
- Licensing options: site license, seat license, single-user.

Deep explanation:

- **Open-source DBMS:** No license fee, but possible costs for support & training.
- **Commercial DBMS:** Paid licenses, but include advanced features, professional support, better scalability.
- **Special modules:**
 - Time-series → financial data
 - Spatial → GIS systems
 - XML → web data exchange
- **Cartridges/blades:** Optional plug-ins for specific domains.

Exam trap: If they ask “What is an Oracle cartridge?” — it’s an **add-on module** for specialized DBMS functionality.

Page 53 — Database Design

Slide content: This likely starts the next module.

Introduction to Database Design Process:

1. **Requirement Analysis** — Understand what the system must do (functional & non-functional requirements).
2. **Conceptual Design** — Build ER diagram capturing entities, attributes, relationships, constraints.
3. **Logical Design** — Convert ER to relational schema, define keys, apply normalization.
4. **Physical Design** — Decide indexing, partitioning, storage formats.
5. **Implementation & Testing** — Create DB, load data, run tests.

Extra exam note:

- Remember: **Conceptual design** ≠ **Logical design**.
- Normalization is part of **logical design**.
- Indexes, partitioning are **physical design** choices.